

Assignment - SQL and R

Jawaid Hakim

2022-09-08

Contents

1	Introduction	1
2	Database schema	2
2.1	Primary Keys	2
2.2	Foreign Keys	2
2.3	Missing Ratings	2
3	Database	2
3.1	Remote AWS (RDS) MySQL Database	2
3.2	Local Database Setup	3
4	Database Connection	4
4.1	Install CNF File	4
4.2	Local Database Connection	4
4.3	Remote Database Connection	4
5	Select Remote or Local Database	5
6	Load FRIENDS and MOVIES	5
7	Load Ratings from Database	6

1 Introduction

SQL databases are widely used repositories for mission critical data. This solution illustrates connection to SQL databases from R. Both local and remote databases are explored. A remote AWS RDS MySQL database is available out of the box. However, a local MySQL database can also be easily set up and configured for use with this solution.

2 Database schema

Schema for this assignment is normalized into 3 tables: FRIENDS, MOVIES, and RATINGS. Normalization has many benefits including efficient storage and reduced operational overhead. Each entity can evolve independently. For instance, adding additional details to FRIENDS, like the date of birth or address, can be done without impacting other entities. Additionally, database maintenance overhead is also reduced - e.g., changing the first name of a FRIEND can be done by updating just a single row.

2.1 Primary Keys

FRIENDS and MOVIES tables each have a primary key named ID. This key is auto-generated by the database and the sequence is initialized to 1 - i.e. first record in table has ID = 1.

2.2 Foreign Keys

Foreign key relationships are created to enforce referential integrity. In this schema, foreign key constraints have been set in the RATINGS table. Click [ER diagram](#) to see foreign key relationships.

2.3 Missing Ratings

It is possible, even highly likely, that not all friends would have viewed all movies. This database schema accommodates this scenario by design - only rated movies need be loaded into RATINGS table.

However, to allow 'NA' ratings to be loaded, the RATINGS.RATING column is defined as an ENUM [0, 1, 2, 3, 4, 5], where 0 (default) is reserved for unrated movies.

3 Database

This solution works with either a pre-configured remote AWS (RDS) MySQL database or a local MySQL database.

3.1 Remote AWS (RDS) MySQL Database

A pre-configured AWS MySQL database has been set up and pre-populated with data to make it easy to run this R solution. Creating the AWS MySQL database was straightforward using the wizard. One caveat was that the default security group rules do not permit access to RDS from external IP addresses. In other words, connection to RDS over the open internet is not allowed.

To get around this limitation, I created a custom security group in AWS and configured its **Inbound rules** as follows: **Protocol:** TCP, **Type:** MYSQL/AURORA, **Port range:** 3306, **Source:** 0.0.0.0/0. Finally, I assigned this security group to the RDS database. RDS can now be accessed over the internet from R Studio and MySQL Workbench.

3.1.1 Test Remote Database Connection

To test connectivity to the RDS MySQL database do the following:

1. Start MySQL Workbench
2. Click + button next to **My SQLConnections**. The **Setup New Connection** windows will launch

3. **Connection name:** Enter any name of your choosing
4. **Hostname:** Enter the **Endpoint** of the RDS MySQL database. Endpoint for the pre-configured RDS database is **cuny-ds.c5iirativieki.us-east-1.rds.amazonaws.com**
5. **Username:** Enter **guest** for the pre-configured RDS database
6. **Password:** Enter **guestpass** for the pre-configured RDS database
7. Click **Test Connection** to validate connection then click **OK** to finish
8. Log into RDS MySQL using the new connection. The **Assignment - SQL and R** schema will be visible. Enjoy!

3.2 Local Database Setup

Make sure you have MySQL and MySQL Workbench installed. Admin database privileges will be required for creating database

3.2.1 Create Local Database

1. Download [schema creation script](#) to local storage
2. In MySQL Workbench click **File - Open SQL Script**, select downloaded file to load into editor, and execute script. On successful execution, schema will be created in MySQL database

3.2.2 Populate Local Database

1. Load data into FRIENDS table
 - Download [FRIENDS data](#) to local storage
 - In MySQL Workbench right-mouse click **Assignment - SQL and R -> Tables -> FRIENDS** table and select **Table Data Import Wizard**. Use wizard to load downloaded FRIENDS data into the FRIENDS table
2. Load data into MOVIES table
 - Download [MOVIES data](#) to local storage
 - In MySQL Workbench right-mouse click **Assignment - SQL and R -> Tables -> MOVIES** table and select **Table Data Import Wizard**. Use wizard to load downloaded MOVIES data into the MOVIES table
3. Load data into RATINGS table
 - Download [RATINGS data](#) to local storage
 - In MySQL Workbench right-mouse click **Assignment - SQL and R -> Tables -> RATINGS** table and select **Table Data Import Wizard**. Use wizard to load downloaded RATINGS data into the RATINGS table

```
library(DBI)
library(RMySQL)
library(RMariaDB)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
library(knitr)
```

4 Database Connection

For security reasons, database connection details are stored in a **CNF** file and not exposed in this R script. This CNF file should be downloaded from github and installed as described below.

4.1 Install CNF File

1. Download [CNF file](#) to local storage in a folder that is accessible by R runtime e.g., current working directory. From the R console run `getwd()` command to see the current working directory. To change the working directory, use the `setwd()` command.

```
cnf.settingsfile <- 'assignment_sql_and_r.cnf'
```

Modify `assignment_sql_and_r.cnf` as appropriate for your MySQL installation. For example, CNF has two configurations: `local_movie_ratings` for a local MySQL database and `remote_movie_ratings` for the pre-configured AWS MySQL. Note: the local configuration assumes there is a `guest` database user.

4.2 Local Database Connection

Connect to the MySQL database using the **MariaDB** driver. Create a function to connect to the local database:

```
cnf.group <- 'local_movie_ratings'  
my.dbConnectLocal <- function() {  
  db <- dbConnect(RMariaDB::MariaDB(), default.file=cnf.settingsfile, group=cnf.group)  
  db  
}
```

4.3 Remote Database Connection

Connect to the AWS RDS MySQL database using the **RMySQL** driver. Create a function to connect to the remote database:

```
cnf.group <- 'remote_movie_ratings'  
my.dbConnectRemote <- function() {  
  db <- dbConnect(RMySQL::MySQL(), default.file=cnf.settingsfile, group=cnf.group)  
  db  
}
```

5 Select Remote or Local Database

Connect to either the remote or local database. Once the connection is made to the appropriate database the rest of the R script remains unchanged:

```
db <- my.dbConnectRemote() # AWS MySQL database
#db <- my.dbConnectLocal() # Local MySQL database
```

For sanity check, list the tables. As expected, FRIENDS, MOVIES, and RATINGS tables are listed.

```
dbListTables(db)
```

```
## [1] "FRIENDS" "MOVIES" "RATINGS"
```

6 Load FRIENDS and MOVIES

Load the FRIENDS table:

```
qry <- 'SELECT * FROM FRIENDS ORDER BY FIRST_NAME, LAST_NAME'
rs <- dbSendQuery(db, qry)
friends <- dbFetch(rs, n=-1)
dbClearResult(rs)
```

```
## [1] TRUE
```

```
head(friends)
```

```
##   ID FIRST_NAME LAST_NAME
## 1  2      Alex  Zakharov
## 2  1   Claudia  Schaab
## 3  4    Howard   Pein
## 4  3      Igor Vaysiberg
## 5  5    Monish   Darda
```

Load the MOVIES table:

```
qry <- 'SELECT * FROM MOVIES ORDER BY TITLE'
rs <- dbSendQuery(db, qry)
movies <- dbFetch(rs, n=-1)
dbClearResult(rs)
```

```
## [1] TRUE
```

```
head(movies)
```

```
##   ID      TITLE
## 1 14  A QUIET PLACE
## 2 13  A STAR IS BORN
## 3 10 BOHEMIAN RHAPSODY
## 4  9          CODA
## 5 11 CRAZY RICH ASIANS
## 6  5    DON'T LOOK UP
```

7 Load Ratings from Database

Finally, load the RATINGS. Since the database schema is normalized, join FRIENDS, MOVIES, and RATINGS tables to load aggregate ratings data.

```
qry <- 'SELECT f.FIRST_NAME, f.LAST_NAME, m.TITLE, r.RATING
        FROM FRIENDS f, MOVIES m, RATINGS r
        WHERE r.FRIEND_ID = f.ID AND r.MOVIE_ID = m.ID
        ORDER BY f.FIRST_NAME, f.LAST_NAME, m.TITLE'
rs <- dbSendQuery(db, qry)
ratings <- dbFetch(rs, n=-1)
dbClearResult(rs)
```

```
## [1] TRUE
```

```
summary(ratings)
```

```
##   FIRST_NAME      LAST_NAME      TITLE      RATING
## Length:70      Length:70      Length:70      Length:70
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
```

```
head(ratings)
```

```
##   FIRST_NAME LAST_NAME      TITLE RATING
## 1      Alex  Zakharov  A QUIET PLACE      3
## 2      Alex  Zakharov  A STAR IS BORN      4
## 3      Alex  Zakharov BOHEMIAN RHAPSODY      4
## 4      Alex  Zakharov          CODA      3
## 5      Alex  Zakharov CRAZY RICH ASIANS      4
## 6      Alex  Zakharov   DON'T LOOK UP      3
```

Enumerations in MySQL are stored as characters. This is not the most convenient representation because we may want to do numeric analysis on ratings, e.g. average movie rating. Convert RATING from character to integer:

```
ratings <- ratings %>% mutate(RATING = as.integer(RATING))
summary(ratings)
```

```
##   FIRST_NAME      LAST_NAME      TITLE      RATING
## Length:70      Length:70      Length:70      Min.   :0.000
## Class :character Class :character Class :character 1st Qu.:3.000
## Mode  :character Mode  :character Mode  :character Median :4.000
##                                     Mean  :3.643
##                                     3rd Qu.:5.000
##                                     Max.   :5.000
```

A rating of 0 is assigned to unrated movies. Filter out data for unrated movies:

```
ratings <- ratings %>% filter(RATING != 0)
summary(ratings)
```

```
##   FIRST_NAME      LAST_NAME      TITLE      RATING
## Length:66      Length:66      Length:66      Min.   :1.000
## Class :character Class :character Class :character 1st Qu.:3.000
## Mode  :character Mode  :character Mode  :character Median :4.000
##                                     Mean  :3.864
##                                     3rd Qu.:5.000
##                                     Max.   :5.000
```

Close (disconnect) database connection:

```
dbDisconnect(db)
```

```
## [1] TRUE
```