

# Project 4

Jawaid Hakim

2022-10-18

## Contents

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr  0.3.4
## v tibble  3.1.8      v dplyr  1.0.10
## v tidyr   1.2.1      v stringr 1.4.1
## v readr   2.1.3      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(readtext)
library(quanteda)
```

```
## Package version: 3.2.3
## Unicode version: 14.0
## ICU version: 70.1
## Parallel computing: 8 of 8 threads used.
## See https://quanteda.io for tutorials and examples.
```

```
library(quanteda.textplots)
library(quanteda.textmodels)
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift
```

```
set.seed(1234)
```

```
sms <- read.delim("SMSSpamCollection.txt",
  sep = '\\t',
  col.names = c('cat', 'sms'),
  quote = "")
```

```
corp_sms <- corpus(sms,
                  text_field = "sms")


# Add numeric document id, used later to partition train/test sets
corp_sms$id_numeric <- 1:ndoc(corp_sms)

# tokenize
toks_sms <- tokens(corp_sms,
                  remove_punct = TRUE,
                  remove_numbers = TRUE,
                  remove_symbols = TRUE,
                  split_hyphens = TRUE) %>%
  tokens_remove(pattern = c(stopwords("en"), "lt", "gt"),
                valuetype = "fixed",
                padding = FALSE,
                min_nchar = 2) %>%
  tokens_tolower(keep_acronyms = TRUE) %>%
  tokens_wordstem()

# doc frequency matrix
dfm_sms <- dfm(toks_sms)
dfm_sms <- dfm_trim(dfm_sms,
                  min_termfreq = 50)
topfeatures(dfm_sms)
```

```
## call  now  go  get  can  ur  just  come  ok  free
## 672 494 451 450 408 391 375 303 285 283
```

```
textplot_wordcloud(dfm_sms,
                  max_words = 50,
                  rotation = 0.1,
                  color = "darkred")
```



```
fcm_sms <- fcm(dfm_sms,
              context = "document",
              count = "frequency",
              window = 5L)
```

```
topfeatures(fcm_sms,
            n = 20,
            scheme = "docfreq")
```


```
##  phone    keep    care    one number  thing  everi    last    come    night    also
##    138     123     118     116    116    115    112    111    110    108    108
##   cant    stop    cos  sleep  start   leav   good   much  person
##    108     106     106     106    105    105    101    101    101
```

```
feat <- names(topfeatures(fcm_sms,
                          n = 20,
                          scheme = "docfreq"))
```

```
fcm_sms_select = fcm_select(fcm_sms,
                            pattern = feat,
                            selection = "keep")
```

```
size <- log(colSums(dfm_select(dfm_sms,
                               feat,
                               selection = "keep"))))
```

```
textplot_network(fcm_sms_select,
  min_freq = 0.8,
  vertex_size = size / max(size) * 3)
```



```
corp_sms_sz <- length(docnames(corp_sms))

# train with 75% data (25% reserved for testing)
id_train <- sample(1:corp_sms_sz, corp_sms_sz * 0.75, replace = FALSE)
head(id_train, 10)
```

```
## [1] 3715 3085 1097 5379 5205 1057 1916 357 998 2015
```

```
dfm_training <- dfm_subset(dfm_sms, id_numeric %in% id_train)
dfm_test <- dfm_subset(dfm_sms, !id_numeric %in% id_train)
```

```
# Naive Bayes classifier for texts
tmod_nb <- textmodel_nb(dfm_training,
  dfm_training$cat,
  prior = "docfreq") # uniform, docfreq, termfreq
summary(tmod_nb)
```

```
##
## Call:
```

```
## textmodel_nb.dfm(x = dfm_training, y = dfm_training$cat, prior = "docfreq")
##
## Class Priors:
## (showing first 2 elements)
##   ham   spam
## 0.8679 0.1321
##
## Estimated Feature Scores:
##           go    great    got    wat    ok    free    win    may
## ham  0.024890 0.006343 0.014532 0.0061020 0.0169410 0.003774 0.001124 0.003051
## spam 0.006862 0.002196 0.001921 0.0002745 0.0008235 0.048861 0.013450 0.002196
##           text    txt    dun    say    already    think    live
## ham  0.004496 0.001044 0.0030510 0.0081895 0.0057005 0.009073 0.001525
## spam 0.029371 0.034038 0.0002745 0.0002745 0.0002745 0.003019 0.006313
##           around    hey    week    now    word    back    like    still
## ham  0.0039342 0.006263 0.004978 0.01895 0.002248 0.007306 0.015335 0.009153
## spam 0.0008235 0.000549 0.019215 0.04008 0.003568 0.005490 0.002745 0.001372
##           send    even    per    friend    custom    prize    claim
## ham  0.008591 0.005299 0.0007226 0.005379 0.0008029 8.029e-05 8.029e-05
## spam 0.014548 0.001647 0.0096075 0.003568 0.0107055 1.812e-02 2.361e-02
```

```
dfm_matched <- dfm_match(dfm_test, features = featnames(dfm_training))
```

```
actual_class <- dfm_matched$cat
predicted_class <- predict(tmod_nb, newdata = dfm_matched)
tab_class <- table(actual_class, predicted_class)
tab_class
```

```
##           predicted_class
## actual_class  ham spam
##           ham  1177  22
##           spam   30 165
```

Sensitivity (True positive rate) = (True Positive)/(True Positive + False Negative) Specificity = (True Negative)/(True Negative + False Positive)

```
confusionMatrix(tab_class, mode = "everything", positive = "spam")
```

```
## Confusion Matrix and Statistics
##
##           predicted_class
## actual_class  ham spam
##           ham  1177  22
##           spam   30 165
##
##           Accuracy : 0.9627
##           95% CI : (0.9514, 0.972)
##           No Information Rate : 0.8659
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.8423
##
```

```
## McNemar's Test P-Value : 0.3317
##
##      Sensitivity : 0.8824
##      Specificity : 0.9751
##      Pos Pred Value : 0.8462
##      Neg Pred Value : 0.9817
##      Precision : 0.8462
##      Recall : 0.8824
##      F1 : 0.8639
##      Prevalence : 0.1341
##      Detection Rate : 0.1184
##      Detection Prevalence : 0.1399
##      Balanced Accuracy : 0.9287
##
##      'Positive' Class : spam
##
```